

```

import tensorflow as tf
from keras.layers import Dense, Flatten, Dropout
from keras.models import Model
from keras.applications.vgg19 import VGG19
from keras.utils import image_dataset_from_directory
import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt

os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

# data pipeline
data = image_dataset_from_directory('data', image_size=(224, 224))

Found 1536 files belonging to 3 classes.


# get a single batch (32 images) within pipeline
data_iterator = data.as_numpy_iterator()
batch = data_iterator.next()

# image shape
print(batch[0].shape)
# y values
print(batch[1])

(32, 224, 224, 3)
[0 2 2 2 2 0 1 0 1 2 2 1 1 1 0 1 2 0 0 1 1 0 1 1 0 1 1 0 0 2 2 0]

# 0 : coquette :: 1 : grunge :: 2 : y2k
fig, ax = plt.subplots(3, 3, figsize = (5, 5))
for i in range(3):
    for j in range(3):
        ax[i, j].imshow(batch[0][i * 3 + j].astype(int))
        ax[i, j].title.set_text(batch[1][i * 3 + j])
        ax[i, j].axis('off')
plt.show()

```



```
# scale data to between 0 and 1 for rgb values instead of 0 to 255
data = data.map(lambda x, y: (x/255, y))
data = data.map(lambda x, y: (x, tf.one_hot(y, depth=3)))
```

```
scaled_iterator = data.as_numpy_iterator()
batch = scaled_iterator.next()
for i in range(5):
    print(batch[1][i])
```

```
[0. 1. 0.]
[0. 0. 1.]
[0. 0. 1.]
[0. 0. 1.]
[1. 0. 0.]
```

```
# since rgb values are now between 0 and 1, don't cast to int
# [1, 0, 0] : coquette :: [0, 1, 0] : grunge :: [0, 0, 1] : y2k
fig, ax = plt.subplots(3, 3, figsize = (5, 5))
```

```
for i in range(3):
    for j in range(3):
        ax[i, j].imshow(batch[0][i * 3 + j])
        ax[i, j].title.set_text(batch[1][i * 3 + j])
        ax[i, j].axis('off')
plt.show()
```

```
/Users/ashley/miniforge3/envs/cv/lib/python3.10/site-packages/
matplotlib/text.py:1279: FutureWarning: elementwise comparison failed;
returning scalar instead, but in the future will perform elementwise
comparison
```

```
if s != self._text:
```



```
num_batches = len(data)
train_size = int((num_batches) * .8)
val_size = int((num_batches) * .1)
test_size = int((num_batches) * .1)
print(str(train_size) + " " + str(val_size) + " " + str(test_size))
```

```
38 4 4
```

```
train = data.take(train_size)
val = data.skip(train_size).take(val_size)
test = data.skip(train_size + val_size).take(test_size)

vgg = VGG19(input_shape = (224, 224, 3), weights = 'imagenet',
include_top = False)
for layer in vgg.layers:
    layer.trainable = False
x = Flatten()(vgg.output)
prediction = Dense(3, activation='softmax')(x)
```

```

model = Model(inputs = vgg.input, outputs = prediction)
model.compile(
    loss='categorical_crossentropy',
    optimizer=tf.keras.optimizers.legacy.Adam(learning_rate=1e-3),
    metrics=['categorical_accuracy']
)
print(model.summary())

```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_1 (Dense)	(None, 3)	75267

```
=====
Total params: 20,099,651
Trainable params: 75,267
Non-trainable params: 20,024,384
```

None

```
from keras.callbacks import EarlyStopping
early_stop = EarlyStopping(monitor = 'val_loss', patience = 2)

results = model.fit(train,
                    validation_data = val,
                    epochs = 20,
                    callbacks = early_stop,
                    batch_size = 32, shuffle = True)
```

Epoch 1/20

```
38/38 [=====] - 41s 956ms/step - loss: 0.9825
- categorical_accuracy: 0.6028 - val_loss: 0.5147 -
val_categorical_accuracy: 0.7500
```

Epoch 2/20

```
38/38 [=====] - 40s 1s/step - loss: 0.3697 -
categorical_accuracy: 0.8758 - val_loss: 0.3622 -
val_categorical_accuracy: 0.8828
```

Epoch 3/20

```
38/38 [=====] - 39s 988ms/step - loss: 0.2392
- categorical_accuracy: 0.9309 - val_loss: 0.3193 -
val_categorical_accuracy: 0.9141
```

Epoch 4/20

```
38/38 [=====] - 40s 1s/step - loss: 0.1749 -
categorical_accuracy: 0.9622 - val_loss: 0.2938 -
val_categorical_accuracy: 0.8828
```

Epoch 5/20

```
38/38 [=====] - 42s 1s/step - loss: 0.1055 -
categorical_accuracy: 0.9877 - val_loss: 0.2708 -
val_categorical_accuracy: 0.9062
```


Epoch 6/20

```
38/38 [=====] - 42s 1s/step - loss: 0.0847 -
categorical_accuracy: 0.9910 - val_loss: 0.2323 -
val_categorical_accuracy: 0.9297
```


Epoch 7/20

```
38/38 [=====] - 44s 1s/step - loss: 0.0658 -  
categorical_accuracy: 0.9951 - val_loss: 0.1566 -  
val_categorical_accuracy: 0.9531  
Epoch 8/20  
38/38 [=====] - 45s 1s/step - loss: 0.0533 -  
categorical_accuracy: 0.9967 - val_loss: 0.1989 -  
val_categorical_accuracy: 0.9375  
Epoch 9/20  
38/38 [=====] - 44s 1s/step - loss: 0.0449 -  
categorical_accuracy: 0.9975 - val_loss: 0.1089 -  
val_categorical_accuracy: 0.9688  
Epoch 10/20  
38/38 [=====] - 45s 1s/step - loss: 0.0415 -  
categorical_accuracy: 0.9951 - val_loss: 0.1228 -  
val_categorical_accuracy: 0.9688  
Epoch 11/20  
38/38 [=====] - 46s 1s/step - loss: 0.0350 -  
categorical_accuracy: 0.9984 - val_loss: 0.0697 -  
val_categorical_accuracy: 0.9922  
Epoch 12/20  
38/38 [=====] - 46s 1s/step - loss: 0.0293 -  
categorical_accuracy: 0.9984 - val_loss: 0.1745 -  
val_categorical_accuracy: 0.9375  
Epoch 13/20  
38/38 [=====] - 47s 1s/step - loss: 0.0276 -  
categorical_accuracy: 0.9984 - val_loss: 0.0549 -  
val_categorical_accuracy: 1.0000  
Epoch 14/20  
38/38 [=====] - 48s 1s/step - loss: 0.0232 -  
categorical_accuracy: 0.9992 - val_loss: 0.1136 -  
val_categorical_accuracy: 0.9688  
Epoch 15/20  
38/38 [=====] - 47s 1s/step - loss: 0.0224 -  
categorical_accuracy: 0.9992 - val_loss: 0.0953 -  
val_categorical_accuracy: 0.9609
```

```
plt.figure(figsize=(4, 2.5))  
plt.plot(results.history['categorical_accuracy'], label = 'train  
accuracy')  
plt.plot(results.history['val_categorical_accuracy'], label =  
'validation accuracy')  
plt.legend()  
plt.show()
```



```
plt.figure(figsize=(4, 2.5))
plt.plot(results.history['loss'], label = 'train loss')
plt.plot(results.history['val_loss'], label = 'val loss')
plt.legend()
plt.show()
```



```
model.evaluate(test, batch_size = 32)
```

```
4/4 [=====] - 7s 709ms/step - loss: 0.0972 - categorical_accuracy: 0.9844
```

```
[0.09721124172210693, 0.984375]
```

```
y_true, y_pred, x_wrong, actual, y_wrong = [], [], [], [], []
```

```
for x, y in test:
    true = tf.argmax(y, axis = 1)
    pred = tf.argmax(model.predict(x), axis = 1)
    for i in range(len(np.array(true))):
        if np.array(true)[i] != np.array(pred)[i]:
            x_wrong.append(np.array(x[i]))
            actual.append(np.array(true)[i])
            y_wrong.append(np.array(pred)[i])
    y_true.append(true)
```



```

    y_pred.append(pred)
y_true = tf.concat(y_true, axis = 0)
y_pred = tf.concat(y_pred, axis=0)

1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 63ms/step

# 0 : coquette :: 1 : grunge :: 2 : y2k
fig, ax = plt.subplots(1, len(x_wrong), figsize = (20, 20))
for i in range(len(x_wrong)):
    ax[i].imshow(x_wrong[i])
    ax[i].title.set_text("Predicted: " + str(y_wrong[i]) + "\nActual: " + str(actual[i]))
    ax[i].axis('off')
plt.show()

```



```

# classification report
from sklearn.metrics import classification_report, confusion_matrix
print(classification_report(y_true, y_pred))

```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	45
1	0.93	0.93	0.93	42
2	0.95	0.95	0.95	41
accuracy			0.95	128
macro avg	0.95	0.95	0.95	128
weighted avg	0.95	0.95	0.95	128

```

print(confusion_matrix(y_true, y_pred, normalize = "true"))
print(confusion_matrix(y_true, y_pred))

```

```

[[0.97777778 0.02222222 0.          ]
 [0.02380952 0.92857143 0.04761905]
 [0.          0.04878049 0.95121951]]
[[44  1  0]
 [ 1 39  2]
 [ 0  2 39]]

```